

Neural Networks & Deep Learning Assignment 07 → ICP 07

Name: Azharuddin Mohammad

700#: 700756035

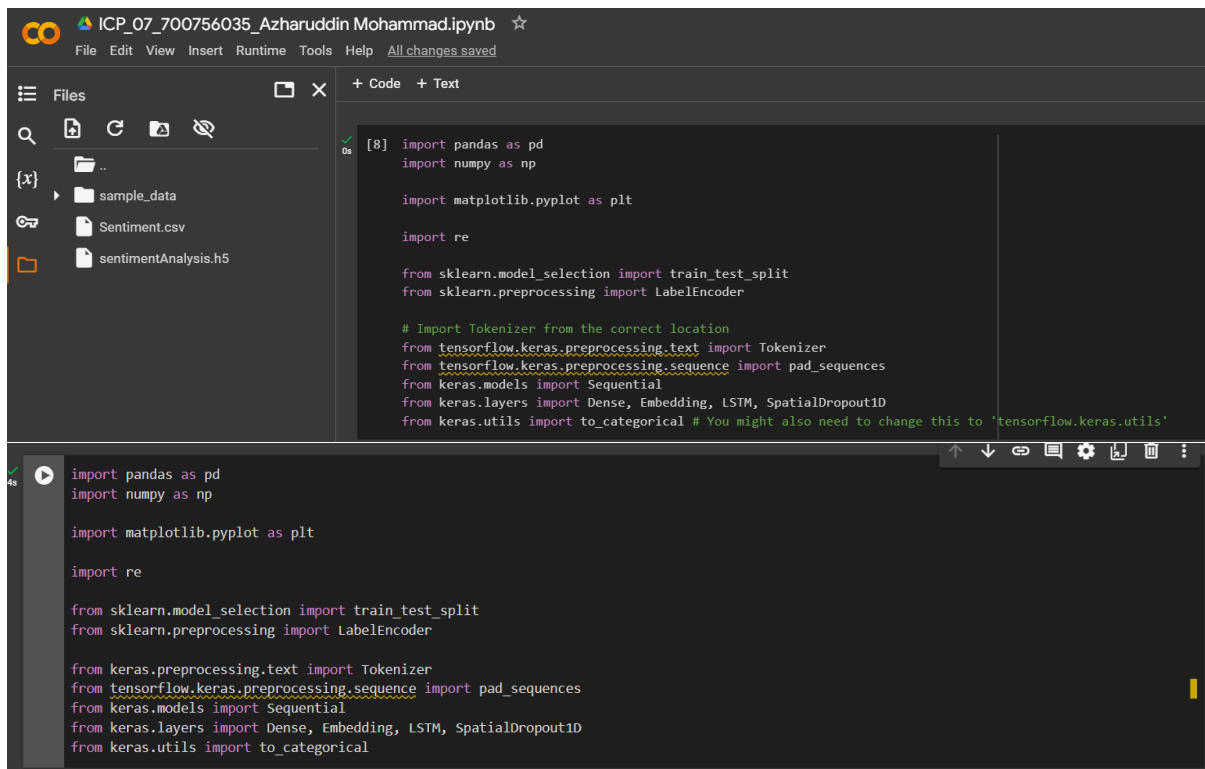
GitHub Link: [https://github.com/mohammadazharuddin982/NN-DL ICP 07 700756035 Azharuddin Mohammad](https://github.com/mohammadazharuddin982/NN-DL_ICP_07_700756035_Azharuddin_Mohammad)

Video Link: <https://vimeo.com/1023119949/006db408a1?share=copy>

Programming elements:

1. Basics of LSTM
2. Types of RNN
3. Use case: Sentiment Analysis on the Twitter data set

Question 01: Save the model and use the saved model to predict on new text data (ex, “A lot of good things are happening. We are respected again throughout the world, and that's a great thing.@realDonaldTrump”)



The screenshot displays a Jupyter Notebook titled "ICP_07_700756035_Azharuddin Mohammad.ipynb". The left sidebar shows a file explorer with a folder named "sample_data" containing "Sentiment.csv" and "sentimentAnalysis.h5". The main area shows two code cells. The top cell, labeled "[8]", contains the following code:

```
[8] import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

import re

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Import Tokenizer from the correct location
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from keras.utils import to_categorical # You might also need to change this to 'tensorflow.keras.utils'
```

The bottom cell contains the following code:

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

import re

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from keras.utils import to_categorical
```

```
ICP_07_700756035_Azharuddin Mohammad.ipynb
File Edit View Insert Runtime Tools Help All changes saved

Files
{ }
sample_data
Sentiment.csv
sentimentAnalysis.h5

+ Code + Text
data = pd.read_csv('Sentiment.csv')
# Keeping only the necessary columns
data = data[['text','sentiment']]

data['text'] = data['text'].apply(lambda x: x.lower())
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-z0-9\s]', '', x))

for idx, row in data.iterrows():
    row[0] = row[0].replace('rt', ' ')

max_fatures = 2000
tokenizer = Tokenizer(num_words=max_fatures, split=' ')
tokenizer.fit_on_texts(data['text'].values)
X = tokenizer.texts_to_sequences(data['text'].values)

X = pad_sequences(X)

embed_dim = 128
lstm_out = 196
def createmodel():
    model = Sequential()
    model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1]))
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
    model.add(Dense(3,activation='softmax'))
    model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy'])
    return model
# print(model.summary())
```

```
+ Code + Text
model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1]))
model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(3,activation='softmax'))
model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy'])
return model
# print(model.summary())

labelencoder = LabelEncoder()
integer_encoded = labelencoder.fit_transform(data['sentiment'])
y = to_categorical(integer_encoded)
X_train, X_test, Y_train, Y_test = train_test_split(X,y, test_size = 0.33, random_state = 42)

batch_size = 32
model = createmodel()
model.fit(X_train, Y_train, epochs = 1, batch_size=batch_size, verbose = 2)
score,acc = model.evaluate(X_test,Y_test,verbose=2,batch_size=batch_size)
print(score)
print(acc)
print(model.metrics_names)

<ipython-input-9-18626b796642>:9: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will
row[0] = row[0].replace('rt', ' ')
<ipython-input-9-18626b796642>:9: FutureWarning: Series.__setitem__ treating keys as positions is deprecated. In a future version, integer keys will
row[0] = row[0].replace('rt', ' ')
291/291 - 50s - loss: 0.8257 - accuracy: 0.6412 - 50s/epoch - 173ms/step
144/144 - 3s - loss: 0.7630 - accuracy: 0.6750 - 3s/epoch - 23ms/step
0.7629973292350769
0.6749672293663025
['loss', 'accuracy']

1m 35s completed at 8:20 PM
```

```
model.save('sentimentAnalysis.h5')

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via `model.save`
saving_api.save_model(

[10] model.save('sentimentAnalysis.h5')

from keras.models import load_model
model= load_model('sentimentAnalysis.h5')
print(integer_encoded)
print(data['sentiment'])

[1 2 1 ... 2 0 2]
0 Neutral
1 Positive
2 Neutral
3 Positive
4 Positive
...
13866 Negative
13867 Positive
13868 Positive
13869 Negative
13870 Positive
Name: sentiment, Length: 13871, dtype: object
```

```
0s sentence = ['A lot of good things are happening. We are respected again throughout the world, and that is a great thing.@realDonaldTrump']
sentence = tokenizer.texts_to_sequences(sentence)
sentence = pad_sequences(sentence, maxlen=28, dtype='int32', value=0)
sentiment_probs = model.predict(sentence, batch_size=1, verbose=2)[0]
sentiment = np.argmax(sentiment_probs)

print(sentiment_probs)
if sentiment == 0:
    print("Neutral")
elif sentiment < 0:
    print("Negative")
elif sentiment > 0:
    print("Positive")
else:
    print("Cannot be determined")

1/1 - 0s - 312ms/epoch - 312ms/step
[0.6675336 0.10805168 0.22441477]
Neutral
```

2. Apply GridSearchCV on the source code provided in the class

```
!pip install keras==2.12.0

Requirement already satisfied: keras==2.12.0 in /usr/local/lib/python3.10/dist-packages (2.12.0)

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=createmodel,verbose=2)
batch_size= [10, 20, 40]
epochs = [1, 2]
param_grid= {'batch_size':batch_size, 'epochs':epochs}
grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result= grid.fit(X_train,Y_train)

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))

Epoch 1/2
186/186 - 33s - loss: 0.8469 - accuracy: 0.6347 - 33s/epoch - 176ms/step
Epoch 2/2
186/186 - 32s - loss: 0.7047 - accuracy: 0.6977 - 32s/epoch - 170ms/step
47/47 - 2s - loss: 0.7497 - accuracy: 0.6815 - 2s/epoch - 39ms/step
Epoch 1/2
186/186 - 32s - loss: 0.8581 - accuracy: 0.6331 - 32s/epoch - 171ms/step
Epoch 2/2
186/186 - 30s - loss: 0.6864 - accuracy: 0.7046 - 30s/epoch - 160ms/step
47/47 - 1s - loss: 0.7475 - accuracy: 0.6825 - 1s/epoch - 26ms/step
Epoch 1/2
186/186 - 34s - loss: 0.8347 - accuracy: 0.6385 - 34s/epoch - 183ms/step
Epoch 2/2
186/186 - 29s - loss: 0.6856 - accuracy: 0.7029 - 29s/epoch - 157ms/step
47/47 - 1s - loss: 0.7845 - accuracy: 0.6733 - 1s/epoch - 28ms/step
Epoch 1/2
465/465 - 67s - loss: 0.8150 - accuracy: 0.6496 - 67s/epoch - 143ms/step
Epoch 2/2
465/465 - 61s - loss: 0.6723 - accuracy: 0.7129 - 61s/epoch - 132ms/step
Best: 0.681911 using {'batch_size': 20, 'epochs': 2}
```

```
ICP_07_700756035_Azharuddin Mohammad.ipynb
File Edit View Insert Runtime Tools Help
Files
sample_data
Sentiment.csv
sentimentAnalysis.h5
+ Code + Text
2. Apply GridSearchCV on the source code provided in the class
import pandas as pd
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

data = pd.read_csv('Sentiment.csv')
target = data['sentiment']
features = data[['sentiment']]

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# fit and transform the target variable
target = label_encoder.fit_transform(target)

# fit and transform the features
features['sentiment'] = label_encoder.fit_transform(features['sentiment'])

model = KerasClassifier(build_fn=createmodel,verbose=2)
batch_size= [10, 20, 40]
epochs = [1, 2]
param_grid= {'batch_size':batch_size, 'epochs':epochs}
grid = GridSearchCV(estimator=model, param_grid=param_grid)
```

```
param_grid= {'batch_size':batch_size, 'epochs':epochs}
grid = GridSearchCV(estimator=model, param_grid=param_grid)

# Split your data into training and testing sets using the selected features
X_train, X_test, Y_train, Y_test = train_test_split(features, target, test_size=0.2, random_state=42)

# Convert features to float32 before fitting
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

grid_result= grid.fit(X_train,Y_train)

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```

<ipython-input-20-91fb039f2583>:29: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
features['sentiment'] = label_encoder.fit_transform(features['sentiment'])

<ipython-input-20-91fb039f2583>:31: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras (<https://github.com/adriangb/scikeras>) instead

```
model = KerasClassifier(build_fn=createmodel,verbose=2)
888/888 - 15s - loss: 0.0401 - accuracy: 0.9991 - 15s/epoch - 17ms/step
222/222 - 1s - loss: 9.0246e-05 - accuracy: 1.0000 - 1s/epoch - 5ms/step
888/888 - 16s - loss: 0.0411 - accuracy: 0.9971 - 16s/epoch - 18ms/step
222/222 - 1s - loss: 8.3490e-05 - accuracy: 1.0000 - 739ms/epoch - 3ms/step
888/888 - 16s - loss: 0.0418 - accuracy: 0.9979 - 16s/epoch - 18ms/step
222/222 - 1s - loss: 9.1541e-05 - accuracy: 1.0000 - 729ms/epoch - 3ms/step
888/888 - 14s - loss: 0.0408 - accuracy: 0.9985 - 14s/epoch - 16ms/step
222/222 - 1s - loss: 8.4153e-05 - accuracy: 1.0000 - 1s/epoch - 5ms/step
888/888 - 17s - loss: 0.0408 - accuracy: 0.9971 - 17s/epoch - 20ms/step
```

✓ 10m 10s completed at 8:53 PM

```
+ Code + Text
```

```
222/222 - 8s - loss: 0.1494 - accuracy: 0.9953 - 8s/epoch - 35ms/step
56/56 - 0s - loss: 9.5393e-04 - accuracy: 1.0000 - 466ms/epoch - 8ms/step
222/222 - 8s - loss: 0.1536 - accuracy: 0.9948 - 8s/epoch - 35ms/step
56/56 - 0s - loss: 0.0010 - accuracy: 1.0000 - 486ms/epoch - 9ms/step
Epoch 1/2
222/222 - 6s - loss: 0.1564 - accuracy: 0.9897 - 6s/epoch - 28ms/step
Epoch 2/2
222/222 - 4s - loss: 6.0396e-04 - accuracy: 1.0000 - 4s/epoch - 16ms/step
56/56 - 1s - loss: 3.0339e-04 - accuracy: 1.0000 - 706ms/epoch - 13ms/step
Epoch 1/2
222/222 - 6s - loss: 0.1566 - accuracy: 0.9861 - 6s/epoch - 28ms/step
Epoch 2/2
222/222 - 3s - loss: 5.9709e-04 - accuracy: 1.0000 - 3s/epoch - 14ms/step
56/56 - 0s - loss: 2.9620e-04 - accuracy: 1.0000 - 467ms/epoch - 8ms/step
Epoch 1/2
222/222 - 6s - loss: 0.1599 - accuracy: 0.9869 - 6s/epoch - 28ms/step
Epoch 2/2
222/222 - 3s - loss: 6.2620e-04 - accuracy: 1.0000 - 3s/epoch - 16ms/step
56/56 - 0s - loss: 3.1835e-04 - accuracy: 1.0000 - 474ms/epoch - 8ms/step
Epoch 1/2
222/222 - 11s - loss: 0.1575 - accuracy: 0.9885 - 11s/epoch - 48ms/step
Epoch 2/2
222/222 - 3s - loss: 5.8291e-04 - accuracy: 1.0000 - 3s/epoch - 14ms/step
56/56 - 1s - loss: 2.8532e-04 - accuracy: 1.0000 - 520ms/epoch - 9ms/step
Epoch 1/2
222/222 - 7s - loss: 0.1541 - accuracy: 0.9973 - 7s/epoch - 29ms/step
Epoch 2/2
222/222 - 5s - loss: 5.7930e-04 - accuracy: 1.0000 - 5s/epoch - 21ms/step
56/56 - 0s - loss: 2.8872e-04 - accuracy: 1.0000 - 467ms/epoch - 8ms/step
1110/1110 - 18s - loss: 0.0345 - accuracy: 0.9978 - 18s/epoch - 16ms/step
Best: 1.000000 using {'batch_size': 10, 'epochs': 1}
```

✓ 10m 10s completed at 8:53 PM